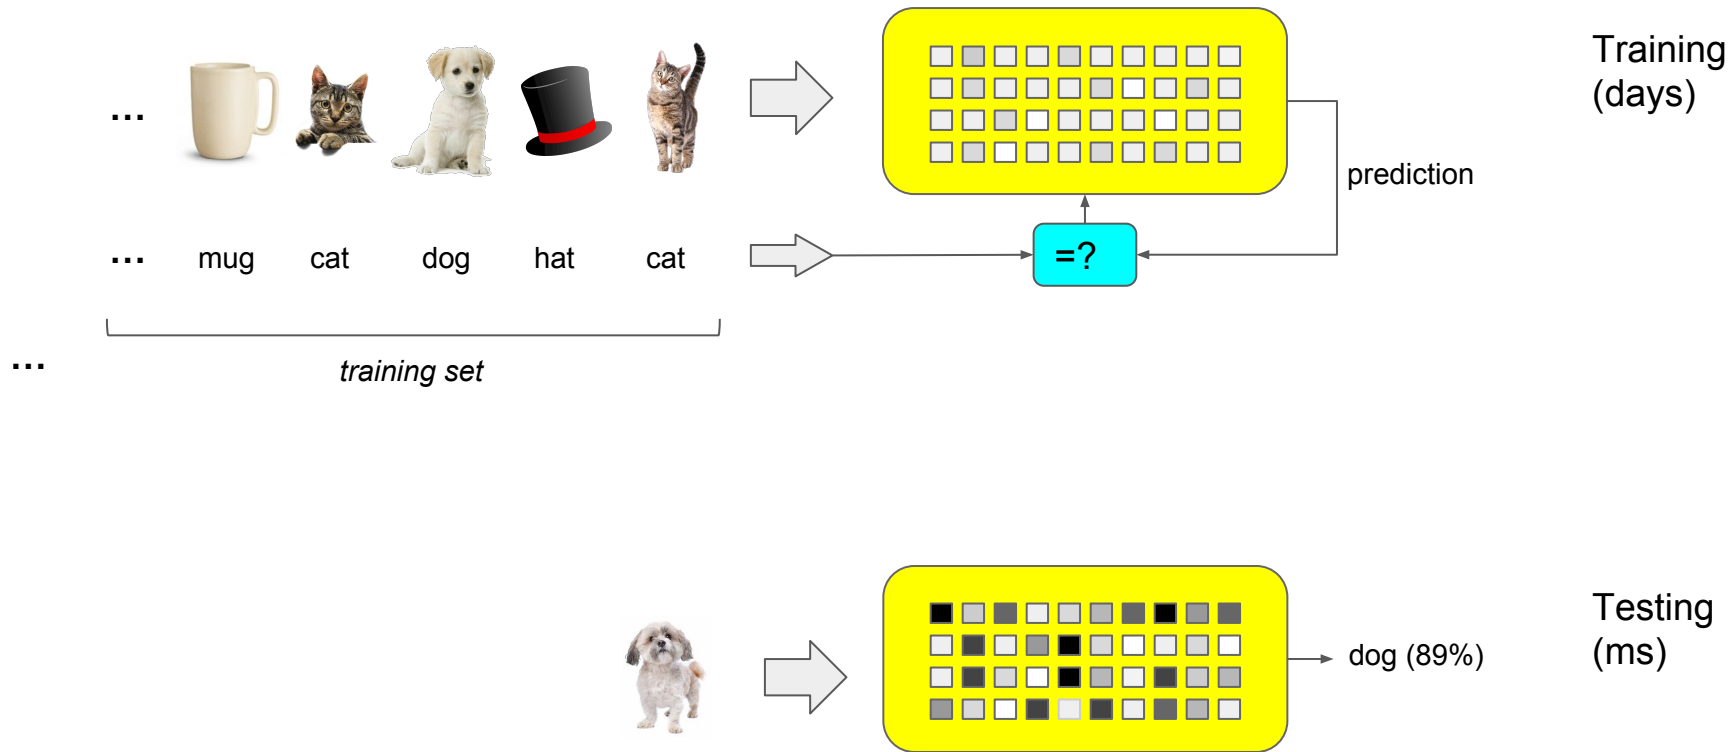


Deep Learning


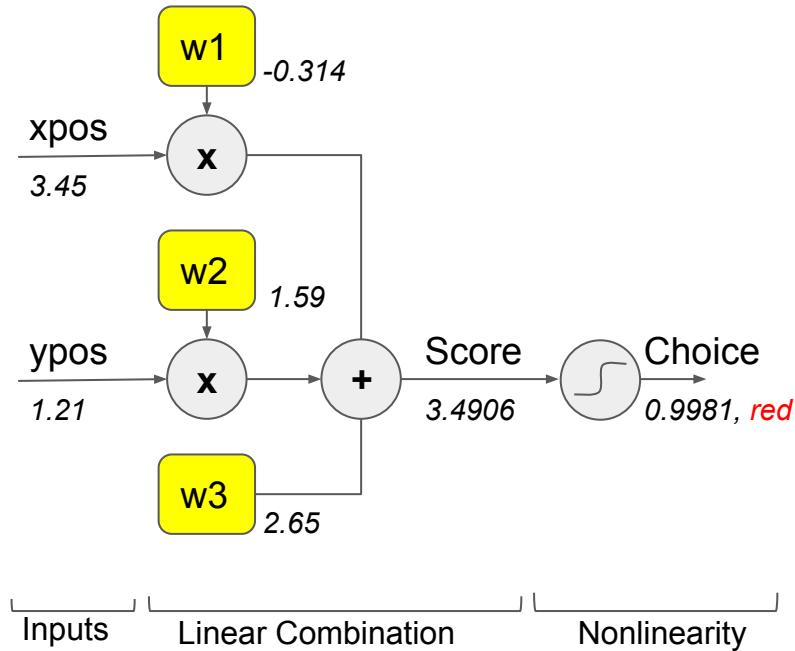
Go deep or go home

Supervised Learning




Single Neuron to Neural Networks

Single Neuron



More Neurons, More Layers: Deep Learning



How to find the best values for the weights?


Define **error** = $|\text{expected} - \text{computed}|^2$

Find parameters that **minimize** average error


Gradient of error wrt w is

$$\frac{\partial \text{error}}{\partial w} = \nabla_w \text{error}$$

Update: $w = w - \text{step} \cdot \nabla_w$




Egg carton in 1 million dimensions



Backpropagation: Assigning Blame

1). For loss to go down by **0.1**, p must go up by **0.05**.
For p to go up by **0.05**, w_1 must go up by **0.09**.
For p to go up by **0.05**, w_2 must go down by **0.12**.
...

3). For p' to go down by **0.01**,
 w_2' must go up by **0.04**



Stochastic Gradient Descent (SGD) and Minibatch

It's too expensive to compute the gradient on all inputs to take a step.

We prefer a quick approximation.

Use a small random sample of inputs (**minibatch**) to compute the gradient.

Apply the gradient to all the weights.

Welcome to SGD, much more efficient than regular GD!

Usually things would get intense at this point...

$$\frac{\partial s}{\partial W_{ij}^{(1)}} = \frac{\partial W_i^{(2)} a_i^{(2)}}{\partial W_{ij}^{(1)}} = \frac{\partial W_i^{(2)} a_i^{(2)}}{\partial W_{ij}^{(1)}} = W_i^{(2)} \frac{\partial a_i^{(2)}}{\partial W_{ij}^{(1)}}$$

Local gradient

$$\Rightarrow W_i^{(2)} \frac{\partial a_i^{(2)}}{\partial W_{ij}^{(1)}} = W_i^{(2)} \frac{\partial a_i^{(2)}}{\partial z_i^{(2)}} \frac{\partial z_i^{(2)}}{\partial W_{ij}^{(1)}}$$

Chain rule

Jacobian matrix

$$= W_i^{(2)} \frac{\sigma'(z_i^{(2)})}{\partial z_i^{(2)}} \frac{\partial z_i^{(2)}}{\partial W_{ij}^{(1)}}$$

Hessian

Sum over paths

$$= W_i^{(2)} \sigma'(z_i^{(2)}) \frac{\partial z_i^{(2)}}{\partial W_{ij}^{(1)}}$$


$$= W_i^{(2)} \sigma'(z_i^{(2)}) \frac{\partial}{\partial W_{ij}^{(1)}} (b_i^{(1)} + a_1^{(1)} W_{i1}^{(1)} + a_2^{(1)} W_{i2}^{(1)} + a_3^{(1)} W_{i3}^{(1)} + a_4^{(1)} W_{i4}^{(1)})$$

$$= W_i^{(2)} \sigma'(z_i^{(2)}) \frac{\partial}{\partial W_{ij}^{(1)}} (b_i^{(1)} + \sum_k a_k^{(1)} W_{ik}^{(1)})$$



$$= W_i^{(2)} \sigma'(z_i^{(2)}) a_j^{(1)}$$

$$= \delta_i^{(2)} \cdot a_j^{(1)}$$

Partial derivatives



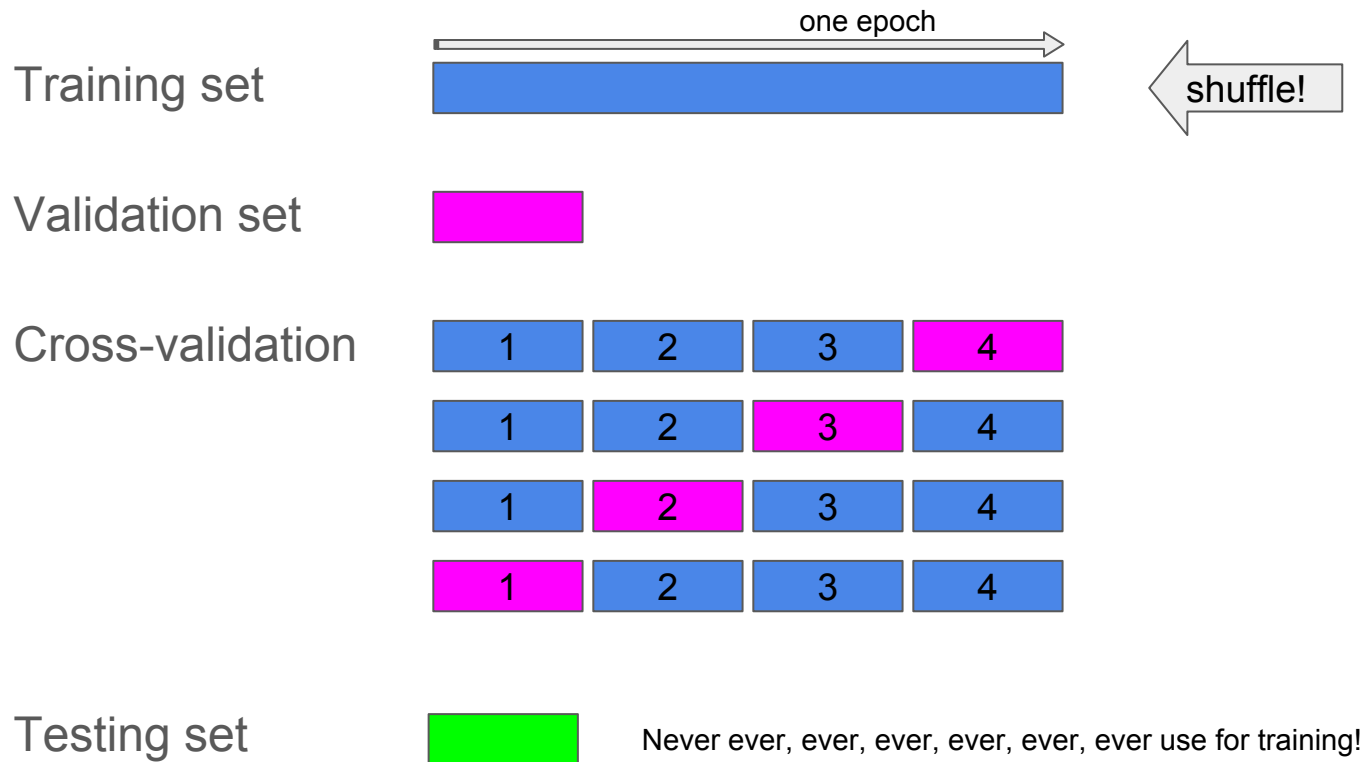
Learning a new representation



Credit: <http://colah.github.io/posts/2014-03-NN-Manifolds-Topology/>

Training Data

Train. Validate. Test.



Tensors

Tensors are not scary

[r=0.45 g=0.84 b=0.76]



Number: 0D

Vector: 1D

Matrix: 2D

Tensor: 3D, 4D, ... array of numbers


This image is a **3264 x 2448 x 3 tensor**



Different types of networks / layers

Fully-Connected Networks



Per layer, every input connects to every neuron.



Recurrent Neural Networks (RNN)

Appropriate when inputs are sequences.


We'll cover in detail when we work on text.



Convolutional Neural Networks (ConvNets)

Appropriate for image tasks, but not limited to image tasks.

We'll cover in detail in the next session.



Hyperparameters

Activations: a zoo of nonlinear functions

Initializations: Distribution of initial weights. Not all zeros.

Optimizers: driving the gradient descent

Objectives: comparing a prediction to the truth



Regularizers: forcing the function we learn to remain “simple”

...and many more


Activations

Nonlinear functions


Sigmoid, tanh, hard tanh and softsign




ReLU (Rectified Linear Unit) and Leaky ReLU




$$\text{ReLU}(x) = \max(x, 0)$$



Softplus and Exponential Linear Unit (ELU)




$$\text{softplus}(x) = \log(1 + e^x)$$



$$f(x) = \begin{cases} x & \text{if } x > 0 \\ \alpha (\exp(x) - 1) & \text{if } x \leq 0 \end{cases}$$

Softmax



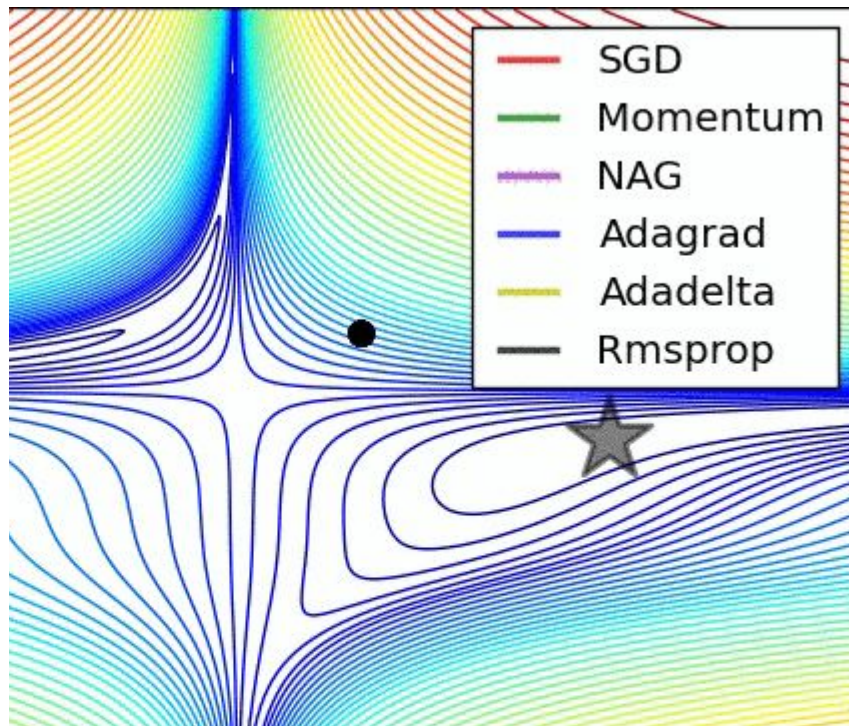

$$p_i = \frac{e^{x_i}}{e^{x_1} + e^{x_2} + e^{x_3} + e^{x_4}}$$

Optimizers

Various algorithms for driving Gradient Descent

Tricks to speed up SGD.

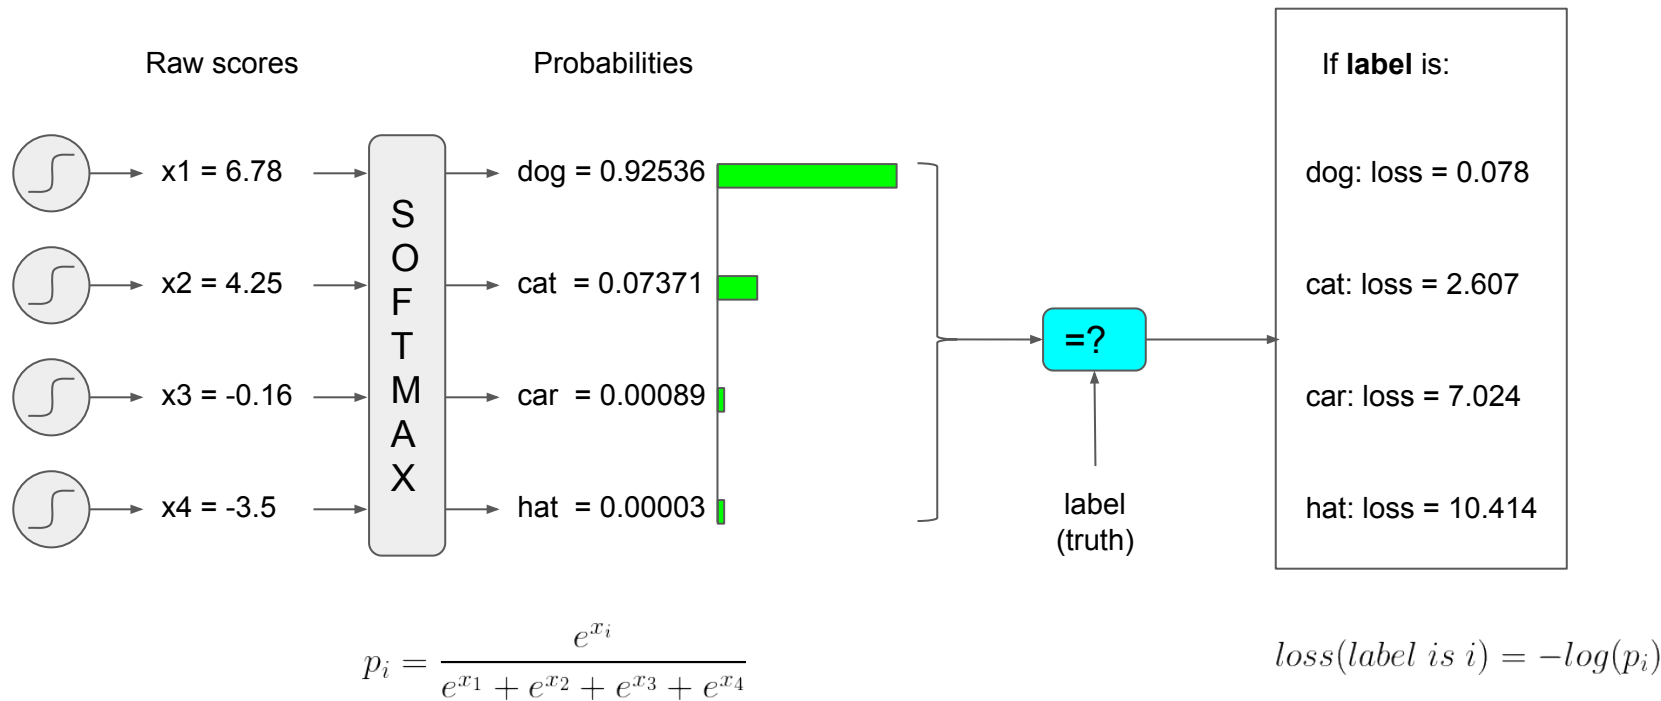
Learn the learning rate.



Cost/Loss/Objective Functions

=?


Cross-entropy Loss



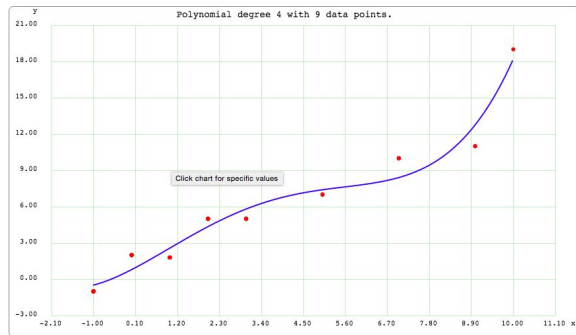
Regularization

Preventing Overfitting

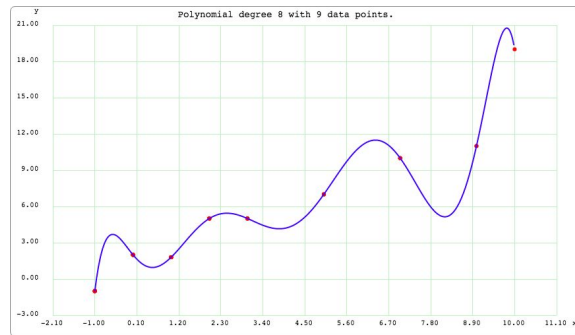
Overfitting



Very simple.
Might not predict well.



Just right?



Overfitting.
Rote memorization.
Will not generalize well.

Regularization: Avoiding Overfitting

Idea: keep the functions “simple” by constraining the weights.

Loss = Error(prediction, truth) + L(keep solution simple)

L2 makes all parameters medium-size $L_2 = \frac{\lambda}{2} \sum_i w_i^2$

L1 kills many parameters.



$$L_1 = \lambda \sum_i |w_i|$$

L1+L2 sometimes used.

Regularization: Dropout


At every step, during training, ignore the output of a fraction p of neurons.

$p=0.5$ is a good default.




One Last Bit of Wisdom

It doesn't work..... why?



It works..... why?



THE END